

Markdown Feature Showcase

A comprehensive example demonstrating everything `md-pdf` can render — typography, code, tables, images, non-ASCII text, emoji, and more.

Table of Contents

- Markdown Feature Showcase
 - Table of Contents
 - 1. Typography & Inline Formatting
 - 2. Emoji & Unicode 🎉
 - Non-ASCII Scripts
 - 3. Images
 - Color Palette
 - Bar Chart
 - Pipeline Diagram
 - 4. Blockquotes
 - 5. Lists
 - Unordered
 - Ordered
 - Task List
 - Definition List
 - 6. Tables
 - Simple Table
 - Aligned Columns
 - 7. Code Samples
 - Python — Async HTTP Client
 - TypeScript — Generic Repository Pattern
 - Rust — Concurrent Task Runner
 - Go — HTTP Middleware Chain
 - SQL — Window Functions & CTEs

- Bash — Deployment Script
 - C — Generic Dynamic Array
 - JSON — Configuration Schema
 - YAML — Kubernetes Deployment
 - HTML + CSS — Card Component
- 8. Footnotes
 - 9. Horizontal Rules
 - 10. Smartsyntax Punctuation
 - 11. Long Paragraph / Prose
-

1. Typography & Inline Formatting

Regular paragraph text with **bold**, *italic*, ~~strikethrough~~, and `inline code`. You can also combine them: ***bold italic*** and ~~**bold strikethrough**~~.

Superscript-like notation: H~2~O and E = mc^2^ (rendered as plain text here).

A [hyperlink to GitHub](#) and a [link with a title](#).

Auto-detected URL: <https://www.python.org>

Email link: hello@example.com

2. Emoji & Unicode



Emoji are fully supported inline: 🚀 🐍 🦀 💎 🌎 🔥 ✅ ✗ ⚠️ 💡 🎨 📄

Non-ASCII Scripts

Greek: Ω φ ε δ — Ἐν ἀρχῇ ἦν ὁ λόγος

Cyrillic: Привет мир! Быстрая коричневая лиса.

Arabic (RTL): مرحباً بالعالم — النص العربي يُكتب من اليمين إلى اليسار:

Chinese (Simplified): 你好，世界！快速的棕色狐狸跳过了懒狗。

Japanese: こんにちは世界！日本語のテキスト。

Korean: 안녕하세요 세계! 빠른 갈색 여우가 게으른 개를 뛰어넘었습니다.

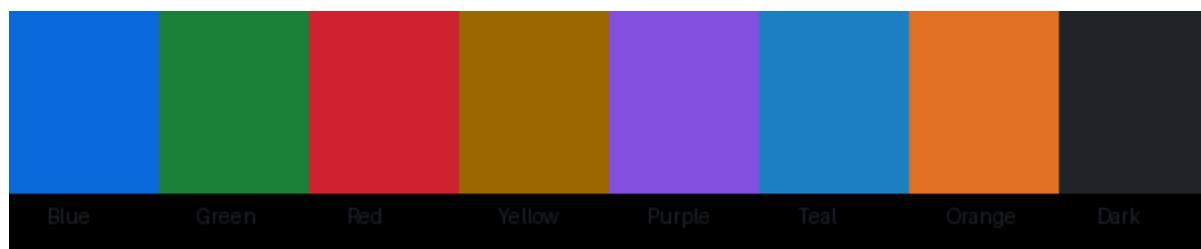
Hebrew: שלום עולם! הטקסט העברי

Mathematical symbols: $\Sigma \int \partial \nabla \infty \approx \neq \leq \geq \in \notin \forall \exists \subset \supset \cap \cup \mathbb{R} \mathbb{C} \mathbb{N} \mathbb{Z}$

Currency: \$ € £ ¥ ₹ ₩ ₧ ₧ ₧

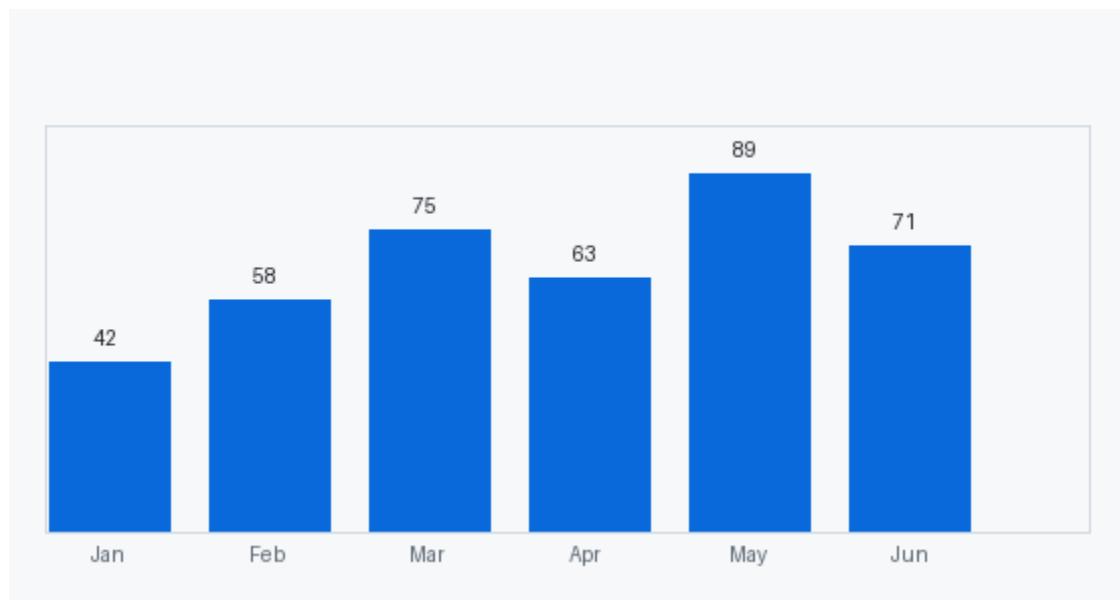
3. Images

Color Palette



Bar Chart

This chart was generated with Pillow to demonstrate image embedding:



Pipeline Diagram



4. Blockquotes

“Programs must be written for people to read, and only incidentally for machines to execute.”

— Harold Abelson, *Structure and Interpretation of Computer Programs*

Nested blockquote:

Outer quote — this is the first level.

Inner quote — this is a nested level, useful for reply threads.

It can span multiple paragraphs.

Back to the outer level.

Blockquote with code inside:

Use `git rebase -i HEAD~3` to interactively rebase the last three commits.
This rewrites history, so **never do this on a shared branch**.

5. Lists

Unordered

- Fruits
- 🍎 Apple

-  Banana
-  Orange
 - Navel orange
 - Blood orange
- Vegetables
-  Broccoli
-  Carrot

Ordered

1. Clone the repository
2. Install dependencies
3. Install Python ≥ 3.13
4. Run `uv sync`
5. Run the application
6. Profit 

Task List

- [x] Design CLI interface
- [x] Add Markdown parser
- [x] Implement syntax highlighting
- [x] Add GitHub-like CSS
- [] Write comprehensive tests
- [] Publish to PyPI

Definition List

Markdown

A lightweight markup language for creating formatted text using a plain-text editor.

PDF

Portable Document Format — a file format developed by Adobe to present documents independent of application software, hardware, and operating systems.

WeasyPrint

A smart solution helping web developers to create PDF documents. It turns simple HTML pages into gorgeous statistical reports, invoices, books and more.

6. Tables

Simple Table

Language	Paradigm	Typing	First appeared
Python	Multi-paradigm	Dynamic	1991
Rust	Systems	Static	2010
JavaScript	Multi-paradigm	Dynamic	1995
Go	Concurrent	Static	2009
Haskell	Functional	Static	1990
SQL	Declarative	Static	1974

Aligned Columns

Metric	Value	Unit	Change
Throughput	142,857	req/s	+12.3%
Latency (p50)	3.2	ms	-0.8ms
Latency (p99)	18.7	ms	-2.1ms
Error rate	0.003	%	+0.001%
Memory usage	512	MB	-64 MB
CPU utilization	34	%	-6%

7. Code Samples

Python — Async HTTP Client

```
import asyncio
import httpx
from dataclasses import dataclass
from typing import Any

@dataclass
class ApiResponse:
    status: int
    data: dict[str, Any]
    headers: dict[str, str]

async def fetch_all(urls: list[str], timeout: float = 10.0) -> list[ApiResponse]:
    """Fetch multiple URLs concurrently and return parsed responses."""
    async with httpx.AsyncClient(timeout=timeout) as client:
        tasks = [client.get(url) for url in urls]
        responses = await asyncio.gather(*tasks, return_exceptions=True)

    results = []
    for resp in responses:
        if isinstance(resp, Exception):
            print(f"Request failed: {resp}")
            continue
        results.append(ApiResponse(
            status=resp.status_code,
            data=resp.json(),
            headers=dict(resp.headers),
        ))
    return results

async def main() -> None:
    urls = [
        "https://api.github.com/repos/python/cpython",
        "https://api.github.com/repos/rust-lang/rust",
    ]
    responses = await fetch_all(urls)
    for r in responses:
        print(f"[{r.status}] stars: {r.data.get('stargazers_count', 'n/a')}")

if __name__ == "__main__":
    asyncio.run(main())
```

TypeScript — Generic Repository Pattern

```
interface Entity {
  id: string;
  createdAt: Date;
  updatedAt: Date;
}

interface Repository<T extends Entity> {
  findById(id: string): Promise<T | null>;
  findAll(filter?: Partial<T>): Promise<T[]>;
  create(data: Omit<T, keyof Entity>): Promise<T>;
  update(id: string, data: Partial<Omit<T, keyof Entity>>): Promise<T>;
  delete(id: string): Promise<void>;
}

class InMemoryRepository<T extends Entity> implements Repository<T> {
  private store = new Map<string, T>();

  async findById(id: string): Promise<T | null> {
    return this.store.get(id) ?? null;
  }

  async findAll(filter?: Partial<T>): Promise<T[]> {
    const all = Array.from(this.store.values());
    if (!filter) return all;
    return all.filter(item =>
      Object.entries(filter).every(([k, v]) => item[k as keyof T] === v)
    );
  }

  async create(data: Omit<T, keyof Entity>): Promise<T> {
    const entity = {
      ...data,
      id: crypto.randomUUID(),
      createdAt: new Date(),
      updatedAt: new Date(),
    } as unknown as T;
    this.store.set(entity.id, entity);
    return entity;
  }

  async update(id: string, data: Partial<Omit<T, keyof Entity>>): Promise<T> {
    const existing = await this.findById(id);
    if (!existing) throw new Error(`Entity ${id} not found`);
    const updated = { ...existing, ...data, updatedAt: new Date() };
    this.store.set(id, updated);
    return updated;
  }

  async delete(id: string): Promise<void> {
    this.store.delete(id);
  }
}
```

Rust — Concurrent Task Runner

```

use std::sync::Arc;
use tokio::sync::{Mutex, Semaphore};
use tokio::task::JoinSet;

#[derive(Debug)]
pub struct TaskRunner {
    concurrency: usize,
    results: Arc<Mutex<Vec<String>>>,
}

impl TaskRunner {
    pub fn new(concurrency: usize) -> Self {
        Self {
            concurrency,
            results: Arc::new(Mutex::new(Vec::new())),
        }
    }

    pub async fn run<F, Fut>(&self, tasks: Vec<F>) -> Vec<String>
    where
        F: Fn() -> Fut + Send + 'static,
        Fut: std::future::Future<Output = String> + Send,
    {
        let semaphore = Arc::new(Semaphore::new(self.concurrency));
        let results = Arc::clone(&self.results);
        let mut join_set = JoinSet::new();

        for task in tasks {
            let permit = Arc::clone(&semaphore);
            let results = Arc::clone(&results);

            join_set.spawn(async move {
                let _permit = permit.acquire().await.unwrap();
                let output = task().await;
                results.lock().await.push(output);
            });
        }

        while let Some(res) = join_set.join_next().await {
            if let Err(e) = res {
                eprintln!("Task panicked: {:?}", e);
            }
        }

        let mut guard = self.results.lock().await;
        std::mem::take(&mut *guard)
    }
}

#[tokio::main]
async fn main() {
    let runner = TaskRunner::new(4);
    let tasks: Vec<_> = (0..10)
        .map(|i| move || async move { format!("task-{:} done") })
        .collect();
}

```

```
let outputs = runner.run(tasks).await;
for o in outputs {
    println!("{}");
}
}
```

Go — HTTP Middleware Chain

```
package main

import (
    "context"
    "log/slog"
    "net/http"
    "time"
)

type Middleware func(http.Handler) http.Handler

func Chain(h http.Handler, middlewares ...Middleware) http.Handler {
    for i := len(middlewares) - 1; i >= 0; i-- {
        h = middlewares[i](h)
    }
    return h
}

func Logger(logger *slog.Logger) Middleware {
    return func(next http.Handler) http.Handler {
        return http.HandlerFunc(func(w http.ResponseWriter, r
*http.Request) {
            start := time.Now()
            rec := &statusRecorder{ResponseWriter: w, status: http.StatusOK}
            next.ServeHTTP(rec, r)
            logger.Info("request",
                "method", r.Method,
                "path", r.URL.Path,
                "status", rec.status,
                "duration", time.Since(start),
            )
        })
    }
}

func Timeout(d time.Duration) Middleware {
    return func(next http.Handler) http.Handler {
        return http.HandlerFunc(func(w http.ResponseWriter, r
*http.Request) {
            ctx, cancel := context.WithTimeout(r.Context(), d)
            defer cancel()
            next.ServeHTTP(w, r.WithContext(ctx))
        })
    }
}

type statusRecorder struct {
    http.ResponseWriter
    status int
}

func (r *statusRecorder) WriteHeader(code int) {
    r.status = code
    r.ResponseWriter.WriteHeader(code)
}
```

```
func main() {
    logger := slog.Default()
    mux := http.NewServeMux()
    mux.HandleFunc("/", func(w http.ResponseWriter, _ *http.Request) {
        w.Write([]byte("Hello, world!\n"))
    })

    handler := Chain(mux,
        Logger(logger),
        Timeout(5*time.Second),
    )
    http.ListenAndServe(":8080", handler)
}
```

SQL — Window Functions & CTEs

```
-- Monthly revenue with running total and rank
WITH monthly_revenue AS (
    SELECT
        DATE_TRUNC('month', o.created_at) AS month,
        p.category,
        SUM(oi.quantity * oi.unit_price) AS revenue,
        COUNT(DISTINCT o.customer_id) AS unique_customers
    FROM orders o
    JOIN order_items oi ON oi.order_id = o.id
    JOIN products p ON p.id = oi.product_id
    WHERE o.status = 'completed'
        AND o.created_at >= NOW() - INTERVAL '12 months'
    GROUP BY 1, 2
),
ranked AS (
    SELECT
        month,
        category,
        revenue,
        unique_customers,
        RANK() OVER (PARTITION BY month ORDER BY revenue DESC) AS revenue_rank,
        SUM(revenue) OVER (
            PARTITION BY category
            ORDER BY month
            ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW
        ) AS cumulative_revenue,
        LAG(revenue, 1) OVER (PARTITION BY category ORDER BY month) AS prev_month_revenue
    FROM monthly_revenue
)
SELECT
    month,
    category,
    revenue,
    unique_customers,
    revenue_rank,
    cumulative_revenue,
    ROUND(
        100.0 * (revenue - prev_month_revenue) / NULLIF(prev_month_revenue, 0),
        2
    ) AS mom_growth_pct
FROM ranked
WHERE revenue_rank <= 5
ORDER BY month DESC, revenue_rank;
```

Bash — Deployment Script

```
#!/usr/bin/env bash
set -euo pipefail

APP_NAME="my-service"
DEPLOY_DIR="/opt/${APP_NAME}"
BACKUP_DIR="/opt/backups/${APP_NAME}"
TIMESTAMP=$(date +%Y%m%d_%H%M%S)

log() { echo "[$(date -u +%FT%TZ)] $*" >&2; }
die() { log "ERROR: $*"; exit 1; }

require_cmd() {
    command -v "$1" &>/dev/null || die "Required command not found: $1"
}

require_cmd docker
require_cmd curl

log "Starting deployment of ${APP_NAME} - ${TIMESTAMP}"

# Health check before deployment
HEALTH=$(curl -sf "http://localhost:8080/health" | jq -r '.status' 2>/dev/null || echo "unreachable")
log "Pre-deploy health: ${HEALTH}"

# Backup current version
if [[ -d "${DEPLOY_DIR}" ]]; then
    log "Backing up current deployment..."
    mkdir -p "${BACKUP_DIR}"
    tar -czf "${BACKUP_DIR}/${TIMESTAMP}.tar.gz" -C "${DEPLOY_DIR}" .
fi

# Pull and restart
log "Pulling latest image..."
docker pull "${APP_NAME}:latest"

log "Restarting service..."
docker compose -f "${DEPLOY_DIR}/docker-compose.yml" up -d --remove-orphans

# Wait for healthy
for i in {1..30}; do
    STATUS=$(curl -sf "http://localhost:8080/health" | jq -r '.status' 2>/dev/null || echo "")
    if [[ "${STATUS}" == "ok" ]]; then
        log "Service healthy after ${i}s."
        exit 0
    fi
    sleep 1
done

die "Service failed to become healthy within 30 seconds. Check logs."
```

C — Generic Dynamic Array

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>

typedef struct {
    void    *data;
    size_t   len;
    size_t   cap;
    size_t   elem_size;
} Vec;

Vec vec_new(size_t elem_size) {
    return (Vec){ .data = NULL, .len = 0, .cap = 0, .elem_size =
elem_size };
}

int vec_push(Vec *v, const void *elem) {
    if (v->len == v->cap) {
        size_t new_cap = v->cap == 0 ? 8 : v->cap * 2;
        void  *new_data = realloc(v->data, new_cap * v->elem_size);
        if (!new_data) return -1;
        v->data = new_data;
        v->cap  = new_cap;
    }
    memcpy((char *)v->data + v->len * v->elem_size, elem, v->elem_size);
    v->len++;
    return 0;
}

void *vec_get(const Vec *v, size_t i) {
    if (i >= v->len) return NULL;
    return (char *)v->data + i * v->elem_size;
}

void vec_free(Vec *v) {
    free(v->data);
    v->data = NULL;
    v->len  = 0;
    v->cap  = 0;
}

int main(void) {
    Vec v = vec_new(sizeof(int));
    for (int i = 0; i < 10; i++) vec_push(&v, &i);

    for (size_t i = 0; i < v.len; i++) {
        printf("%d ", *(int *)vec_get(&v, i));
    }
    printf("\n");

    vec_free(&v);
    return 0;
}

```

JSON — Configuration Schema

```
{  
  "$schema": "https://json-schema.org/draft/2020-12/schema",  
  "title": "AppConfig",  
  "description": "Application configuration schema",  
  "type": "object",  
  "required": ["server", "database"],  
  "properties": {  
    "server": {  
      "type": "object",  
      "properties": {  
        "host": { "type": "string", "default": "0.0.0.0" },  
        "port": { "type": "integer", "minimum": 1, "maximum": 65535,  
"default": 8080 },  
        "workers": { "type": "integer", "minimum": 1, "default": 4 },  
        "tls": {  
          "type": "object",  
          "properties": {  
            "enabled": { "type": "boolean" },  
            "cert_file": { "type": "string" },  
            "key_file": { "type": "string" }  
          }  
        }  
      }  
    },  
    "database": {  
      "type": "object",  
      "required": ["url"],  
      "properties": {  
        "url": { "type": "string", "format": "uri" },  
        "pool_size": { "type": "integer", "default": 10 },  
        "pool_timeout": { "type": "number", "default": 30.0 },  
        "echo_sql": { "type": "boolean", "default": false }  
      }  
    },  
    "logging": {  
      "type": "object",  
      "properties": {  
        "level": { "type": "string", "enum": ["DEBUG", "INFO", "WARNING",  
"ERROR"] },  
        "format": { "type": "string", "enum": ["json", "text"] }  
      }  
    }  
  }  
}
```

YAML — Kubernetes Deployment

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: web-api
  namespace: production
  labels:
    app: web-api
    version: "2.4.1"
spec:
  replicas: 3
  selector:
    matchLabels:
      app: web-api
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 0
  template:
    metadata:
      labels:
        app: web-api
  spec:
    containers:
      - name: api
        image: my-org/web-api:2.4.1
        ports:
          - containerPort: 8080
        env:
          - name: DATABASE_URL
            valueFrom:
              secretKeyRef:
                name: db-secret
                key: url
          - name: LOG_LEVEL
            value: "INFO"
        resources:
          requests:
            cpu: "100m"
            memory: "128Mi"
          limits:
            cpu: "500m"
            memory: "512Mi"
    livenessProbe:
      httpGet:
        path: /health
        port: 8080
        initialDelaySeconds: 10
        periodSeconds: 15
    readinessProbe:
      httpGet:
        path: /ready
        port: 8080
        initialDelaySeconds: 5
        periodSeconds: 5
```

HTML + CSS — Card Component

```
<!DOCTYPE html>
<html lang="en">
<head>
  <meta charset="UTF-8">
  <style>
    .card {
      display: flex;
      flex-direction: column;
      gap: 12px;
      padding: 20px;
      border-radius: 12px;
      border: 1px solid #d1d9e0;
      background: #fff;
      max-width: 360px;
      font-family: system-ui, sans-serif;
    }

    .card__avatar {
      width: 48px;
      height: 48px;
      border-radius: 50%;
      background: linear-gradient(135deg, #0969da, #8250df);
    }

    .card__title {
      font-size: 1.1rem;
      font-weight: 600;
      color: #1f2328;
      margin: 0;
    }

    .card__badge {
      display: inline-block;
      padding: 2px 8px;
      border-radius: 999px;
      font-size: 0.75rem;
      background: #ddf4ff;
      color: #0969da;
    }
  </style>
</head>
<body>
  <div class="card">
    <div class="card__avatar"></div>
    <h2 class="card__title">Jane Smith</h2>
    <span class="card__badge">Staff Engineer</span>
    <p>Building reliable distributed systems. Open source contributor. 🦀</p>
  </div>
</body>
</html>
```

8. Footnotes

Markdown was created by John Gruber¹ in 2004 and has since become the de facto standard for writing on the web.²

WeasyPrint converts HTML and CSS to PDF using the Pango³ layout engine.

9. Horizontal Rules

Three ways to draw a rule:

10. Smarty pants Punctuation

The “smarty” extension converts straight quotes and dashes:

- “double quotes” become curly “double quotes”
 - ‘single quotes’ become curly ‘single quotes’
 - – becomes an en dash —
 - — becomes an em dash —
 - ... becomes an ellipsis...
-

11. Long Paragraph / Prose

The history of markup languages stretches back decades. Before the web, documents were typeset using systems like TeX — Donald Knuth’s masterwork, published in 1978 — or troff, the document formatting system developed at Bell Labs. These systems gave authors precise control over layout but required deep expertise to use effectively.

HTML changed everything. By separating *structure* from *presentation*, it allowed non-experts to create documents that could be viewed on any device. But HTML's verbosity made it poorly suited for writing prose. Enter Markdown: a lightweight syntax that compiles to HTML, letting authors focus on content rather than tags.

Today, Markdown powers GitHub READMEs, technical documentation, static site generators, note-taking apps, and — thanks to tools like this one — printable PDFs. The simplicity that made it successful also means there is no single canonical Markdown; dozens of flavors exist, each adding features the original spec omitted: tables, task lists, footnotes, fenced code blocks, strikethrough, and more.

“The overriding design goal for Markdown’s formatting syntax is to make it as readable as possible.”

— John Gruber

Generated by `md-pdf` — <https://github.com> 

1. John Gruber — <https://daringfireball.net> 
2. See also the CommonMark specification at <https://commonmark.org> 
3. Pango is a library for laying out and rendering text — <https://pango.gnome.org> 